



Concept Definition: The Agent Orchestrator

👤 创建者	👤 Yilun Sun
🕒 创建时间	January 22, 2026 9:10 PM
☰ Category	product management
👤 Last editor	👤 Yilun Sun
🕒 Update time	February 21, 2026 10:42 AM
⚙️ Status	Not Started

Concept Definition: The Agent Orchestrator

"Human behavior = A chaotic, unstable, timeout-prone execution environment full of self-deception. Lumi is the Agent Orchestrator built for this 'Human System'."

You can present this comparison to show the 1:1 mapping between Coding Agents and Lumi:

1. The Isomorphism: Coding Agent vs. Habit Agent

Feature	Coding Agent	Habit Agent (Lumi)
Goal	Ship features / Pass tests	Ship sleep / Complete routines
Process	Writing code is just a step	Reminders are just a step
The Hard Part	Reliable execution (no hallucinations/loops)	Reliable execution (no procrastination/self-deception)

Tools	Editor, Terminal, Browser	Notifications, Screen Locks, Timers, Checklists
Failure Modes	Infinite loops, Test failures	Time blindness, Emotional avoidance, "Doomscrolling"
Fix Loop	Plan → Act → Test → Reflect → Patch	Plan → Act → Verify → Reflect → Patch Strategy

Key Insight:

"For a Coding Agent, 'action' is a tool call. For Lumi, 'action' is invoking a human behavior—the human behavior is the API."

2. Core Philosophy: "Human-as-Tool"

Use this to explain "Agency" (AI 主动性):

- **The Concept:** "To an LLM, a tool isn't just a software API; it can be a human behavior."
- **The Implementation:** `tool.sleep_routine.start()` isn't executed by code—it's implemented by guiding the user to brush their teeth, dim the lights, and get into bed.
- **The Definition of Agency:** "Agency isn't the AI sleeping for you. Agency is the AI making the right intervention at the right time to 'invoke' the user's action."
- **Pitch Line:** "If tools are APIs for machines, behaviors are APIs for humans."

3. The Execution Stack

(To prove this is a complex system, not a chatbot)

- **Planner:** Not just "Sleep at 11pm." It schedules the protocol: *22:15 Wrap-up → 22:30 Wash → 22:50 Bed.*
- **Executor:** Pushes the user to the next state via notifications, screen locks, or companion dialogue.
- **Verifier:** Checks if the user actually transitioned states (Signals: Phone idle, charging status, or user confirmation).
- **Memory:** Tracks failure patterns (e.g., "User claims 10 mins but takes 45").

- **Self-Correct:** Handles pushback. (User: "I'm not tired." Agent: "Okay, let's switch to the 5-minute 'combat shower' protocol.")
-

4. Addressing Skepticism (The "Why not an alarm?" Slide)

Don't let them think it's just a reminder app. Frame it as **System Engineering for Humans**:

- "We are applying **agentic software engineering principles** to human execution."
 - "The user is not just the end-user; the user is a component of the **execution environment**."
 - **Differentiation:**
 - Not just Prompts → **State Machines**
 - Not just Chat → **Verifiers**
 - Not just Reminders → **Self-Correction Loops**
-

5. The "One-Liner" Demo Script

"At 10:15 PM, Lumi starts a wrap-up protocol, negotiates deadlines with the user, verifies progress, and self-corrects until the user is in bed—then uses next-day consequence feedback to update future strategies."

6. Product Positioning Statements

Option A (For Tech Judges/Hackathon):

"Lumi is an orchestrator for human execution. It applies the same plan–execute–verify–self-correct loop found in coding agents, but here, the 'tools' are human behaviors: wrap-up, hygiene routines, and sleep."

Option B (Product/User Focused):

"Lumi doesn't just 'remind you to sleep.' It runs a bedtime protocol with you. It negotiates, adapts, and follows through—turning your intention into action."

Option C (Slogan):

|"An agent for your actions, not just your chats."
